

MATLAB 6.5 Image Processing Toolbox Tutorial

The purpose of this tutorial is to gain familiarity with MATLAB's Image Processing Toolbox. This tutorial does not contain all of the functions available in MATLAB. It is very useful to go to **Help\MATLAB Help** in the MATLAB window if you have any questions not answered by this tutorial. Many of the examples in this tutorial are modified versions of MATLAB's help examples. The help tool is especially useful in image processing applications, since there are numerous filter examples.

1. Opening MATLAB in the microcomputer lab

1.1. Access the **Start Menu**, Proceed to **Programs**, Select **MATLAB 6.5** from the **MATLAB 6.5** folder

--OR--

1.2. Open through **C:\MATLAB6p5\bin\win32\matlab.exe**

2. MATLAB

2.1. When MATLAB opens, the screen should look something like what is pictured in Figure 2.1, below.

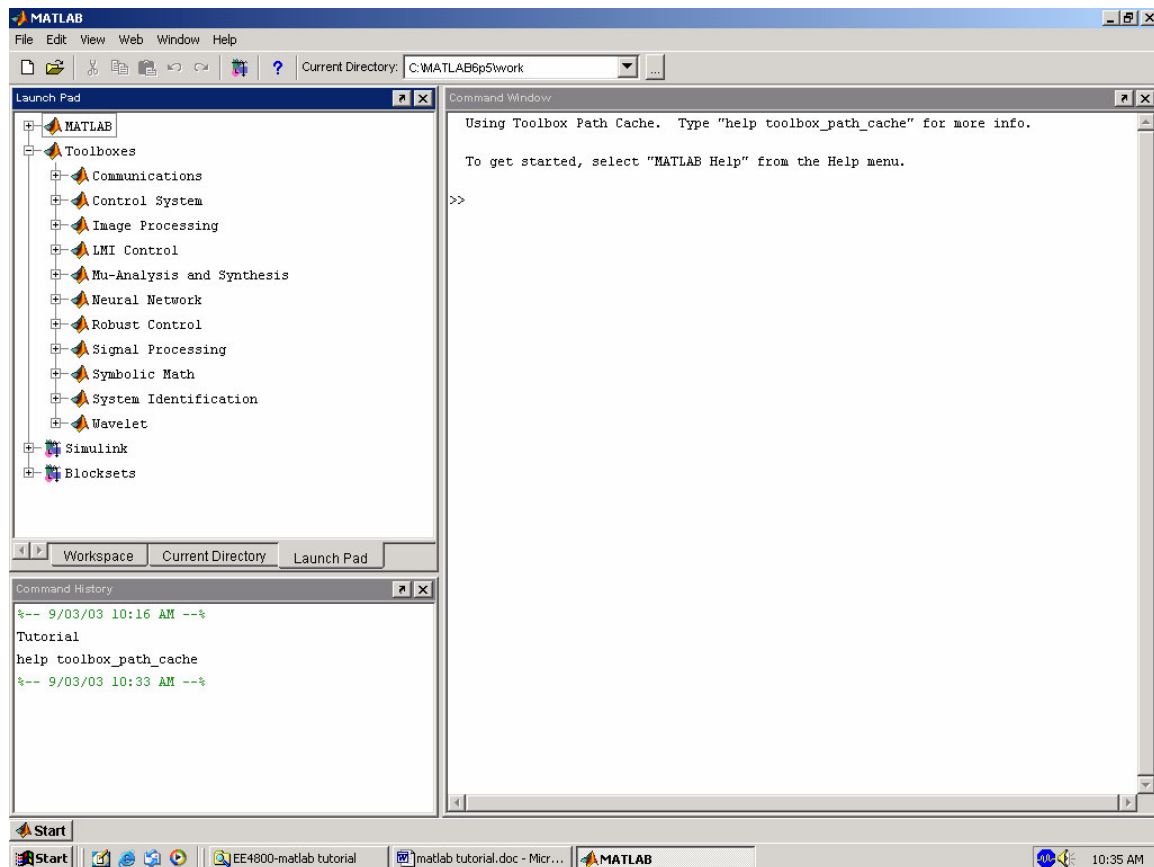


Figure 2.1: MATLAB window

- 2.2. The **Command Window** is the window on the right hand side of the screen. This window is used to both enter commands for MATLAB to execute, and to view the results of these commands.
- 2.3. The **Command History** window, in the lower left side of the screen, displays the commands that have been recently entered into the Command Window.
- 2.4. In the upper left hand side of the screen there is a window that can contain three different windows with tabs to select between them. The first window is the **Current Directory**, which tells the user which M-files are currently in use. The second window is the **Workspace** window, which displays which variables are currently being used and how big they are. The third window is the **Launch Pad** window, which is especially important since it contains easy access to the available toolboxes, of which, Image Processing is one. If these three windows do not all appear as tabs below the window space, simply go to **View** and select the ones you want to appear.
- 2.5. In order to gain some familiarity with the Command Window, try Example 2.1, below. You must type code after the `>>` prompt and press return to receive a new prompt. If you write code that you do not want to reappear in the MATLAB Command Window, you must place a semi colon after the line of code. If there is no semi colon, then the code will print in the command window just under where you typed it.

Example 2.1

```
>> X = 1;           %press enter to go to next line
>> Y = 1;           %press enter to go to next line
>> Z = X + Y         %press enter to receive result
```

As you probably noticed, MATLAB gave an answer of $Z = 2$ under the last line of typed code. If there had been a semi colon after the last statement, the answer would not have been printed. Also, notice how the variables you used are listed in the Workspace Window and the commands you entered are listed in the Command History window. If you want to retype a command, an easy way to do this is to press the \uparrow or \downarrow arrows until you reach the command you want to reenter.

3. The M-file

- 3.1. **M-file** – An M-file is a MATLAB document the user creates to store the code they write for their specific application. **Creating an M-file is highly recommended**, although not entirely necessary. An M-file is useful because it saves the code the user has written for their application. It can be manipulated and tested until it meets the user's specifications. The advantage of using an M-file is that the user, after modifying their code, must only tell MATLAB to run the M-file, rather than reenter each line of code individually.
- 3.2. **Creating an M-file** – To create an M-file, select **File\New ► M-file**.
- 3.3. **Saving** – The next step is to save the newly created M-file. In the M-file window, select **File\Save As...** Choose a location that suits your needs, such as a disk, the hard drive or the U drive. It is not recommended that you work from

your disk or from the U drive, so before editing and testing your M-file you may want to move your file to the hard drive.

- 3.4. **Opening an M-file** – To open up a previously designed M-file, simply open MATLAB in the same manner as described before. Then, open the M-file by going to **File\Open...**, and selecting your file. Then, in order for MATLAB to recognize where your M-file is stored, you must go to **File\Set Path...** This will open up a window that will enable you to tell MATLAB where your M-file is stored. Click the **Add Folder...** button, then browse to find the folder that your M-file is located in, and press **OK**. Then in the Set Path window, select **Save**, and then **Close**. If you do not set the path, MATLAB may open a window saying your file is not in the current directory. In order to get by this, select the “**Add directory to the top of the MATLAB path**” button, and hit **OK**. This is essentially the same as setting the path, as described above.
- 3.5. **Writing Code** – After creating and saving your M-file, the next step is to begin writing code. A suggested first move is to begin by writing comments at the top of the M-file with a description of what the code is for, who designed it, when it was created, and when it was last modified. Comments are declared by placing a % symbol before them. Comments appear in green in the M-file window. See Figure 3.1, below, for Example 3.1.

Example 3.1.

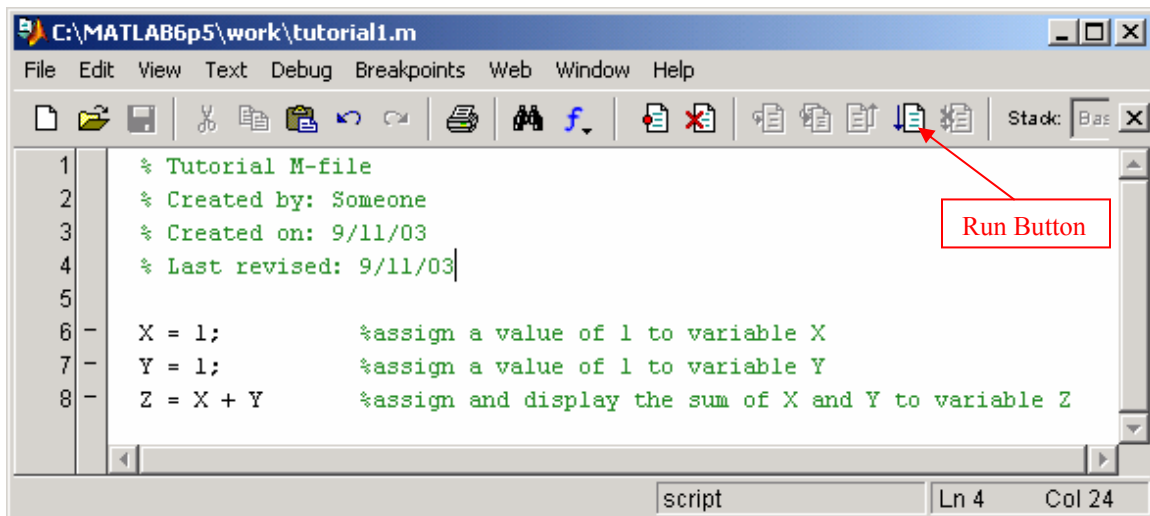


Figure 3.1: Example of M-file

- 3.6. **Resaving** – After writing code, you must save your work before you can run it. Save your code by going to **File\Save**.
- 3.7. **Running Code** – To run code, simply go to the main MATLAB window and type the name of your M-file after the >> prompt. Other ways to run the M-file are to press **F5** while the M-file window is open, select **Debug\Run**, or press the **Run** button (see Figure 3.1) in the M-file window toolbar.

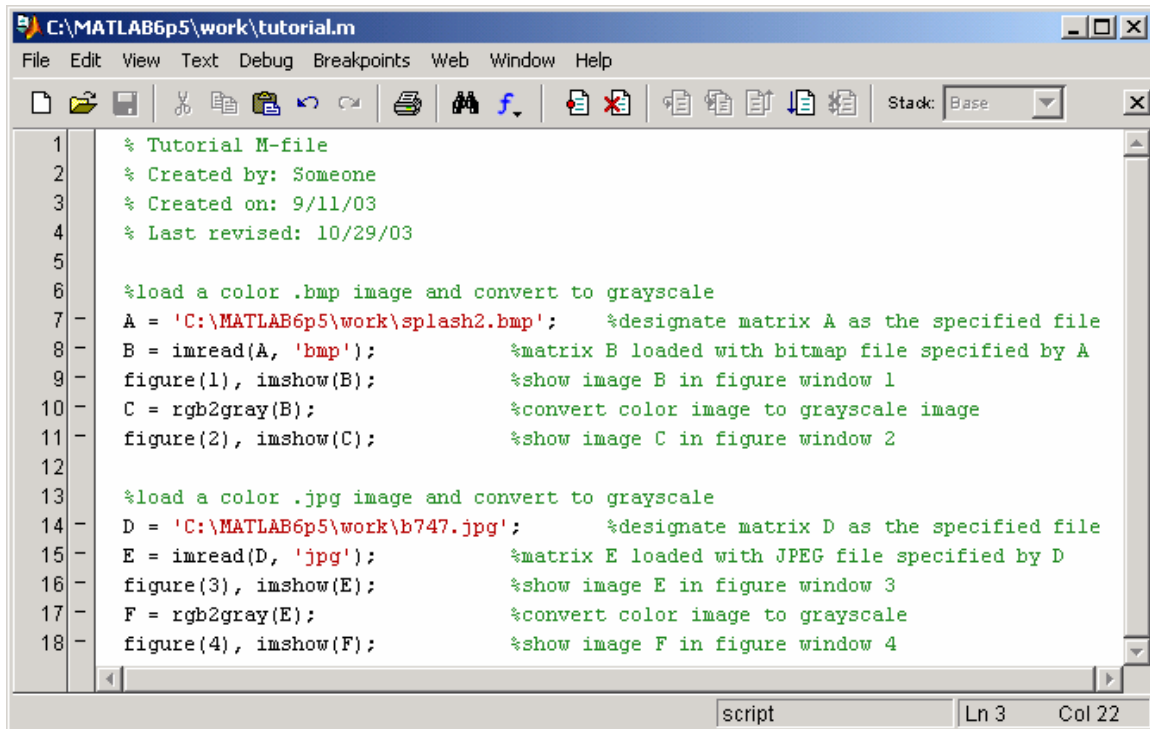
4. Images

- 4.1. Images** – The first step in MATLAB image processing is to understand that a digital image is composed of a two or three dimensional matrix of pixels. Individual pixels contain a number or numbers representing what grayscale or color value is assigned to it. Color pictures generally contain three times as much data as grayscale pictures, depending on what color representation scheme is used. Therefore, color pictures take three times as much computational power to process. In this tutorial the method for conversion from color to grayscale will be demonstrated and all processing will be done on grayscale images. However, in order to understand how image processing works, we will begin by analyzing simple two dimensional 8-bit matrices.
- 4.2. Loading an Image** – Many times you will want to process a specific image, other times you may just want to test a filter on an arbitrary matrix. If you choose to do this in MATLAB you will need to load the image so you can begin processing. If the image that you have is in color, but color is not important for the current application, then you can change the image to grayscale. This makes processing much simpler since then there are only a third of the pixel values present in the new image. Color may not be important in an image when you are trying to locate a specific object that has good contrast with its surroundings. Example 4.1, below, demonstrates how to load different images.

Example 4.1.

In some instances, the image in question is a matrix of pixel values. For example, you may need something to test a filter on, but you do not yet need a real image to test the filter. Therefore, you can simply create a matrix that has the characteristics wanted, such as areas of high and low frequency. See Example 6.1, for a demonstration of this. Other times a stored image must be imported into MATLAB to be processed. If color is not an important aspect then **rgb2gray** can be used to change a color image into a grayscale image. The class of the new image is the same as that of the color image. As you can see from the example M-file in Figure 4.1, MATLAB has the capability of loading many different image formats, two of which are shown. The function **imread** is used to read an image file with a specified format. Consult **imread** in MATLAB's help to find which formats are supported. The function **imshow** displays an image, while **figure** tells MATLAB which figure window the image should appear in. If **figure** does not have a number associated with it, then figures will appear chronologically as they appear in the M-file. Figures 4.2, 4.3, 4.4 and 4.5, below, are a loaded bitmap file, the image in Figure 4.2 converted to a grayscale image, a loaded JPEG file, and the image in Figure 4.4 converted to a grayscale image, respectively. The images used in this example are both MATLAB example images. In order to demonstrate how to load an image file, these images were copied and pasted into the folder denoted in the M-file in Figure 4.1. In Example 7.1, later in this tutorial, you will see that MATLAB images can be loaded by simply using the **imread** function. However, this function will only load an image stored in:

C:\MATLAB6p5\toolbox\images\imdemos. Therefore, it is a good idea to know how to load any image from any folder.



```
1 % Tutorial M-file
2 % Created by: Someone
3 % Created on: 9/11/03
4 % Last revised: 10/29/03
5
6 %load a color .bmp image and convert to grayscale
7 A = 'C:\MATLAB6p5\work\splash2.bmp'; %designate matrix A as the specified file
8 B = imread(A, 'bmp'); %matrix B loaded with bitmap file specified by A
9 figure(1), imshow(B); %show image B in figure window 1
10 C = rgb2gray(B); %convert color image to grayscale image
11 figure(2), imshow(C); %show image C in figure window 2
12
13 %load a color .jpg image and convert to grayscale
14 D = 'C:\MATLAB6p5\work\b747.jpg'; %designate matrix D as the specified file
15 E = imread(D, 'jpg'); %matrix E loaded with JPEG file specified by D
16 figure(3), imshow(E); %show image E in figure window 3
17 F = rgb2gray(E); %convert color image to grayscale
18 figure(4), imshow(F); %show image F in figure window 4
```

script Ln 3 Col 22

Figure 4.1: M-file for Loading Images

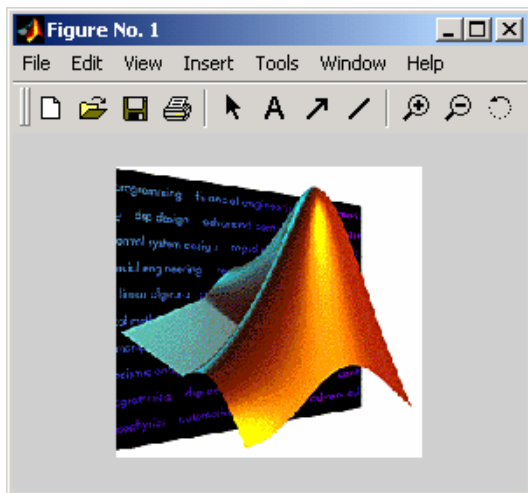


Figure 4.2: Bitmap Image

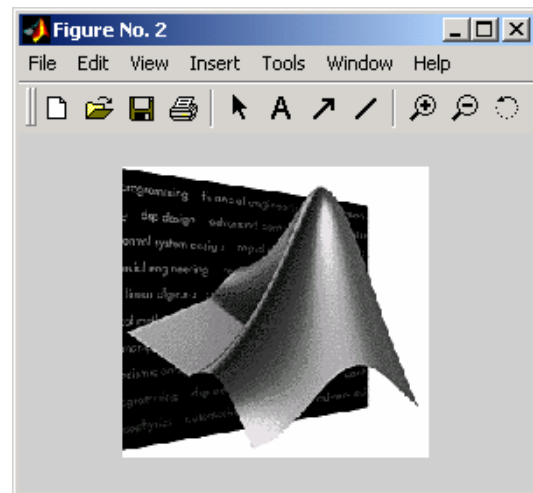


Figure 4.3: Grayscale Image

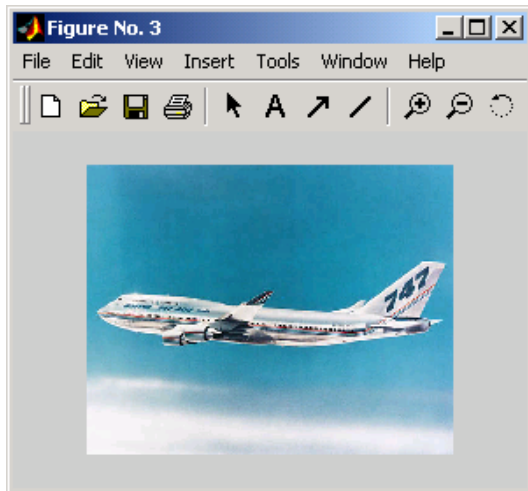


Figure 4.4: JPEG Image

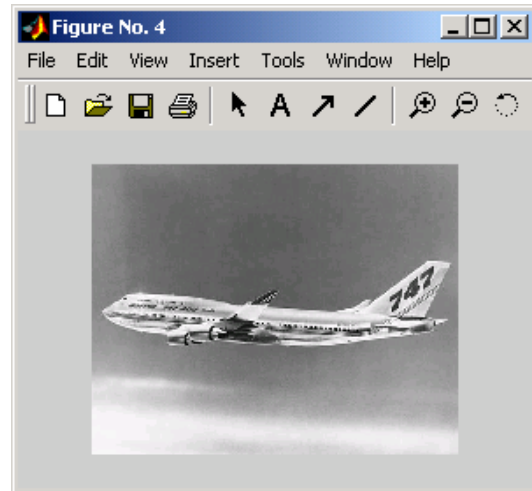


Figure 4.5: Grayscale Image

- 4.3 Writing an Image** – Sometimes an image must be saved so that it can be transferred to a disk or opened with another program. In this case you will want to do the opposite of loading an image, reading it, and instead write it to a file. This can be accomplished in MATLAB using the **imwrite** function. This function allows you to save an image as any type of file supported by MATLAB, which are the same as supported by **imread**. Example 4.2, below, contains code necessary for writing an image.

Example 4.2

In order to save an image you must use the **imwrite** function in MATLAB. The M-file in Figure 4.6 contains code for saving an image. This M-file loads the same bitmap file as described in the M-file pictured in Figure 4.1. However, this new M-file saves the grayscale image created as a JPEG image. Just like in Example 4.1, the “splash2” bitmap picture must be moved into MATLAB’s work folder in order for the **imread** function to find it. When you run this M-file notice how the JPEG image that was created is saved into the work folder.

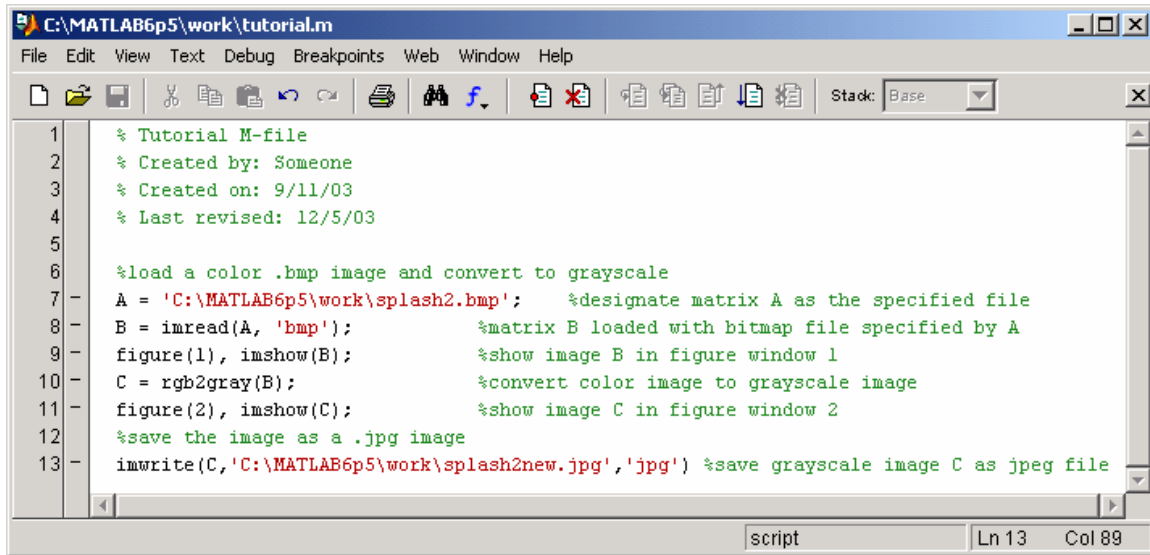


Figure 4.6: M-file for Saving an Image

5. Image Properties

- 5.1. **Histogram** – A histogram is bar graph that shows a distribution of data. In image processing histograms are used to show how many of each pixel value are present in an image. Histograms can be very useful in determining which pixel values are important in an image. From this data you can manipulate an image to meet your specifications. Data from a histogram can aid you in contrast enhancement and thresholding. In order to create a histogram from an image, use the **imhist** function. Contrast enhancement can be performed by the **histeq** function, while thresholding can be performed by using the **graythresh** function and the **im2bw** function. See Example 5.1, for a demonstration of **imhist**, **imadjust**, **graythresh**, and **im2bw**. If you want to see the resulting histogram of a contrast enhanced image, simply perform the **imhist** operation on the image created with **histeq**.
- 5.2. **Negative** – The negative of an image means the output image is the reversal of the input image. In the case of an 8-bit image, the pixels with a value of 0 take on a new value of 255, while the pixels with a value of 255 take on a new value of 0. All the pixel values in between take on similarly reversed new values. The new image appears as the opposite of the original. The **imadjust** function performs this operation. See Example 5.1 for an example of how to use **imadjust** to create the negative of the image. Another method for creating the negative of an image is to use **imcomplement**, which is described in Example 7.5.

Example 5.1

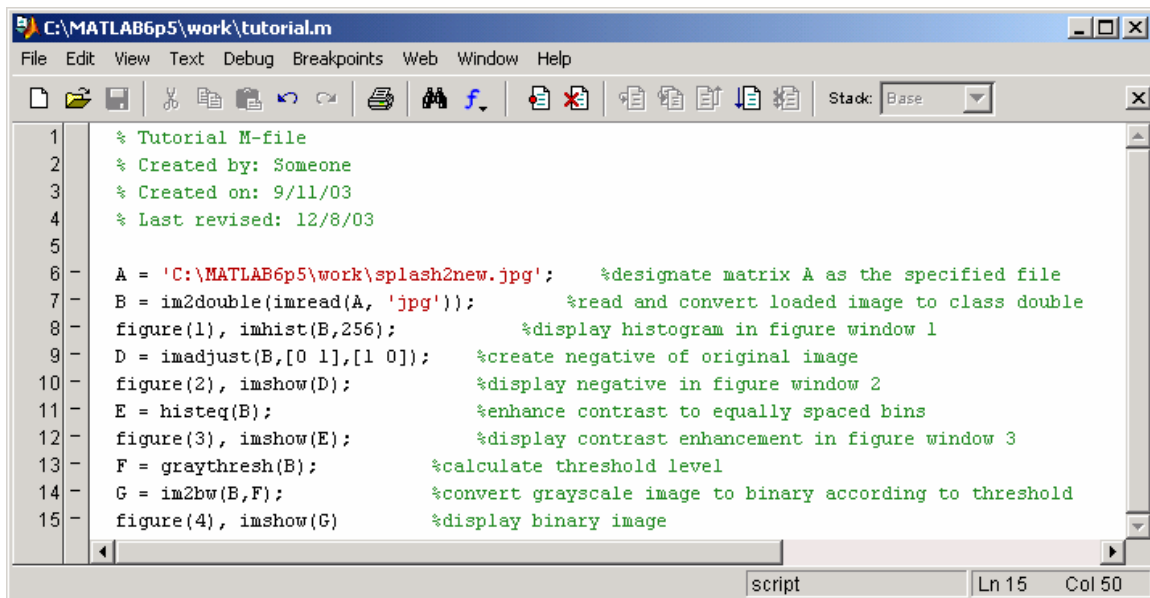
In this example the JPEG image created in Example 4.2 was used to create a histogram of the pixel value distribution and a negative of the original image. The contrast was then enhanced and finally the image was transformed into a binary image according to a certain threshold value. Figure 5.1, below, contains the M-file used to perform these operation. Figure 5.2 contains the histogram of

the image pictured in Figure 4.3. As you can see the histogram gives a distribution between 0 and 1. In order to find the exact pixel value, you must scale the histogram by the number of bits representing each pixel value. In this case, this is an 8-bit image, so scale by 255. As you can see from the histogram, there is a lot of black and white in the image. Figure 5.3 contains the negative of the image pictured in Figure 4.3. Pixel values have been rotated about the midpoint in the histogram. Figure 5.4 contains a contrast enhanced version of the image in Figure 4.3. As you can see, there is some blurring around the edges of the object in the center of the image. However, it is slightly easier to read the words in the image. This is an example of the trade-offs that are common in image processing. In this case, sacrificing fine edges allowed us to see the words better. Figure 5.5 contains a binary image of the image in Figure 4.3. This particular binary image was created according to the threshold level, thresh. The value for thresh was displayed in the MATLAB Command Window as:

```
>>
thresh =
```

```
0.5020
```

MATLAB chooses a value for thresh that minimizes the intraclass variance of black and white pixels. If this value does not meet your expectations, use a different value when using the **im2bw** function. Another function new to this example was **im2double**. This function converts the image from its current class to class double. Many MATLAB functions cannot perform operations on class unit8 or unit16, so they must first be converted into class double. This is due to the unsigned nature of class unit. Certain mathematical functions must be able to output to a floating point array in order to operate. When writing an image, MATLAB converts the data back to class unit.



```

1 % Tutorial M-file
2 % Created by: Someone
3 % Created on: 9/11/03
4 % Last revised: 12/8/03
5
6 A = 'C:\MATLAB6p5\work\splash2new.jpg'; %designate matrix A as the specified file
7 B = im2double(imread(A, 'jpg')); %read and convert loaded image to class double
8 figure(1), imhist(B,256); %display histogram in figure window 1
9 D = imadjust(B,[0 1],[1 0]); %create negative of original image
10 figure(2), imshow(D); %display negative in figure window 2
11 E = histeq(B); %enhance contrast to equally spaced bins
12 figure(3), imshow(E); %display contrast enhancement in figure window 3
13 F = graythresh(B); %calculate threshold level
14 G = im2bw(B,F); %convert grayscale image to binary according to threshold
15 figure(4), imshow(G) %display binary image

```

script Ln 15 Col 50

Figure 5.1: M-file for Creating Histogram, Negative, Contrast Enhanced and Binary Images from the Image Created in Example 4.2

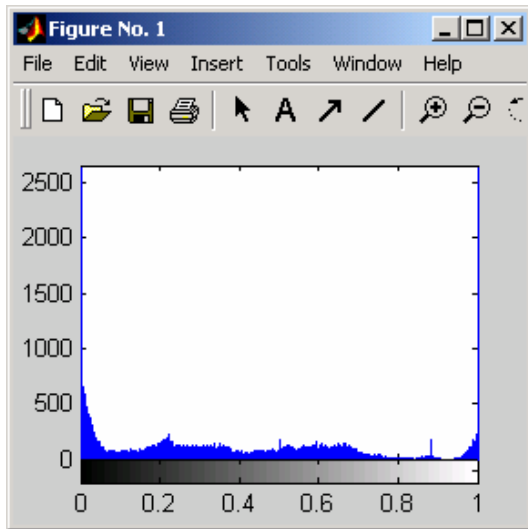


Figure 5.2: Histogram

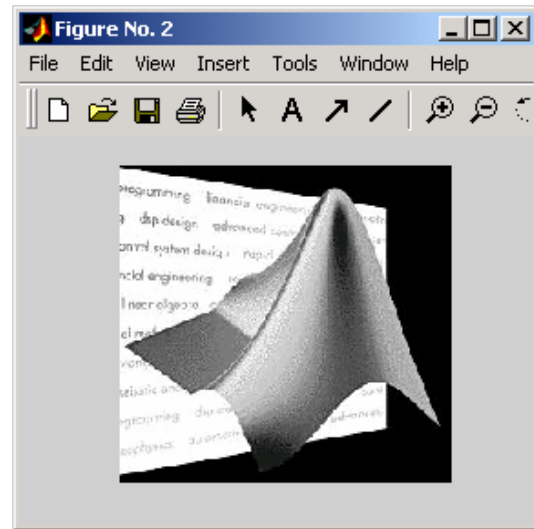


Figure 5.3: Negative

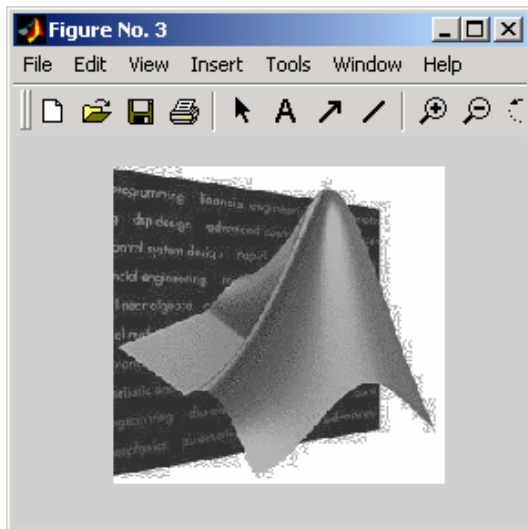


Figure 5.4: Contrast Enhanced

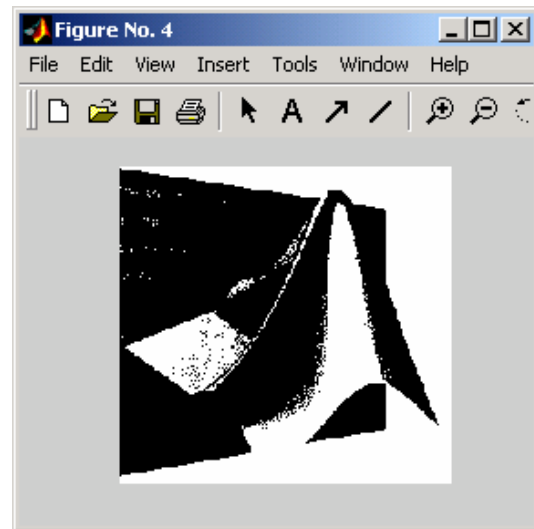


Figure 5.5: Binary

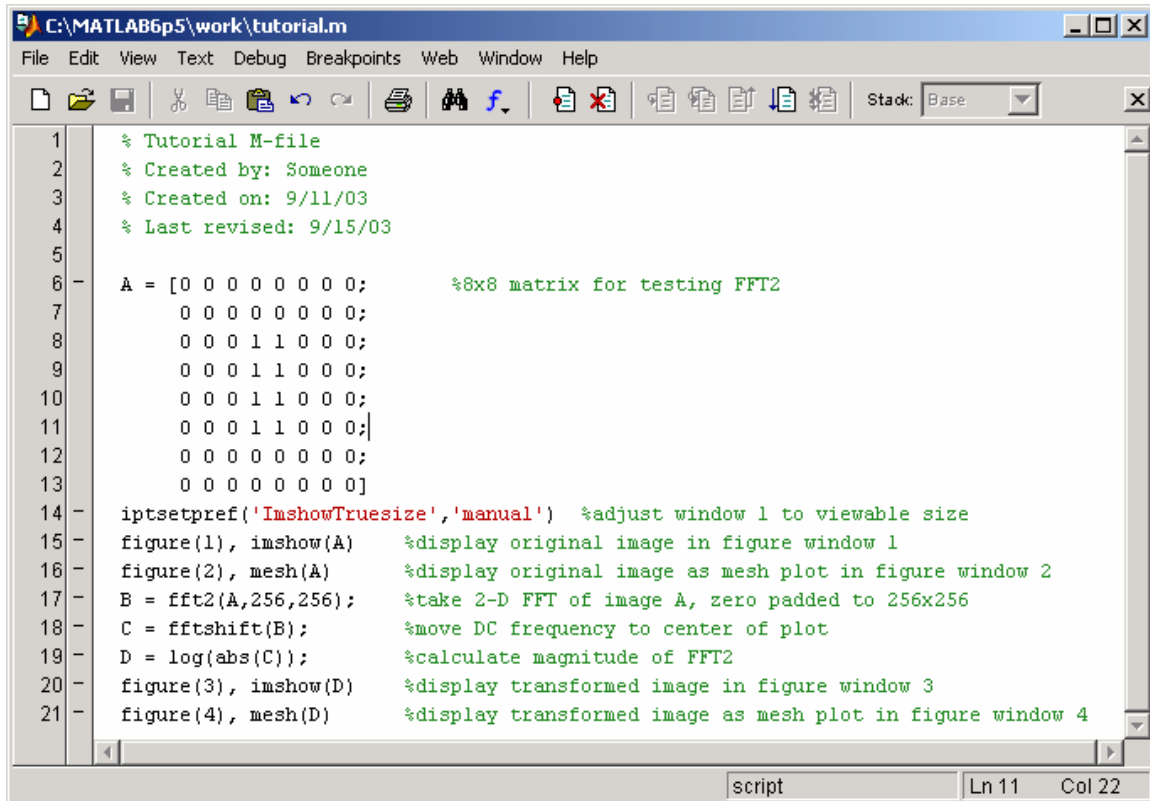
6. Frequency Domain

6.1. Fourier Transform – In order to understand how different image processing filters work, it is a good idea to begin by understanding what frequency has to do with images. An image is in essence a two dimensional collection of discrete signals. Therefore, the signals have frequencies associated with them. For instance, if there is relatively little change in grayscale values as you scan across an image, then there is lower frequency content contained within the image. If there is wide variation in grayscale values across an image then there will be more frequency content associated with the image. This may seem somewhat confusing, so let us think about this in terms that are more familiar to us. From

signal processing, we know that any signal can be represented by a collection of sine waves of differing frequencies, magnitudes and phases. This transformation of a signal into its constituent sinusoids is known as the Fourier Transform. This collection of sine waves can potentially be infinite, if the signal is difficult to represent, but is generally truncated at a point where adding more signals does not significantly improve the resolution of the recreation of the original signal. In digital systems, we use a Fourier Transform designed in such a way that we can enter discrete input values, specify our sampling rate, and have the computer generate discrete outputs. This is known as the Discrete Fourier Transform, or DFT. MATLAB uses a fast algorithm for performing a DFT, which is called the Fast Fourier Transform, or FFT, whose MATLAB command is **fft**. The FFT can be performed in two dimensions, **fft2** in MATLAB. This is very useful in image processing because we can then determine the frequency content of an image. Still confused? Picture an image as a two dimensional matrix of signals. If you plotted just one row, so that it showed the grayscale value stored within each pixel, you might end up with something that looks like a bar graph, with varying values in each pixel location. Each pixel value in this signal may appear to have no correlation to the next one. However, the Fourier Transform can determine which frequencies are present in the signal. In order to see the frequency content, it is useful to view the absolute value of the magnitude of the Fourier Transform, since the output of a Fourier Transform is complex in nature. See Example 6.1, below, for a demonstration of how to perform a two dimensional FFT on an image.

Example 6.1

In this example, we will construct an 8x8 test matrix, *A*, and perform a two dimensional Fast Fourier Transform on it. The M-file used to do this is pictured in Figure 6.1, below. When viewed, the original image is a white rectangle on a black background, as shown in Figure 6.2. In MATLAB, black is denoted as 0, while white is the highest number in the matrix. In this case white is 1. When 8 bits are used to represent grayscale, white is 255. Figure 6.3, below, is the mesh plot of the original image pictured in Figure 6.2. Mesh plots are created using the **mesh** function.



```
1 % Tutorial M-file
2 % Created by: Someone
3 % Created on: 9/11/03
4 % Last revised: 9/15/03
5
6 A = [0 0 0 0 0 0 0 0;      %8x8 matrix for testing FFT2
7      0 0 0 0 0 0 0 0;
8      0 0 0 1 1 0 0 0;
9      0 0 0 1 1 0 0 0;
10     0 0 0 1 1 0 0 0;
11     0 0 0 1 1 0 0 0;
12     0 0 0 0 0 0 0 0;
13     0 0 0 0 0 0 0 0];
14
15 iptsetpref('imshowTrueSize','manual') %adjust window 1 to viewable size
16 figure(1), imshow(A) %display original image in figure window 1
17 figure(2), mesh(A) %display original image as mesh plot in figure window 2
18 B = fft2(A,256,256); %take 2-D FFT of image A, zero padded to 256x256
19 C = fftshift(B); %move DC frequency to center of plot
20 D = log(abs(C)); %calculate magnitude of FFT2
21 figure(3), imshow(D) %display transformed image in figure window 3
22 figure(4), mesh(D) %display transformed image as mesh plot in figure window 4
```

Figure 6.1: M-File for Fourier Transform

Figure 6.4, below, is the image of the two dimensional FFT of the image in Figure 6.2. As you can see, Figure 6.4 is quite different from Figure 6.2. Figure 6.2 is a representation of the matrix's pixel values in space, while Figure 6.4 is a representation of which frequencies are present within the matrix (the 0, DC, frequency is in the center). When moving from left to right across the center of the image in Figure 6.2, you encounter a short pulse, which requires many more sine terms to represent it than the wide pulse you encounter as you move vertically across the image in Figure 6.2. This is evident in Figure 6.4. As you can see, as you move from left to right across the image, you encounter more instances of frequencies being present in the original image. As you move vertically, you do not encounter as many instances of frequencies being present. A shorter pulse requires more frequency components to represent it. Figure 6.5, below, is the mesh plot of the image in Figure 6.4.

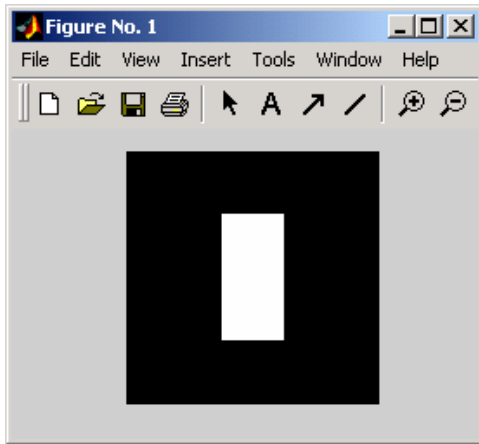


Figure 6.2: Original Image

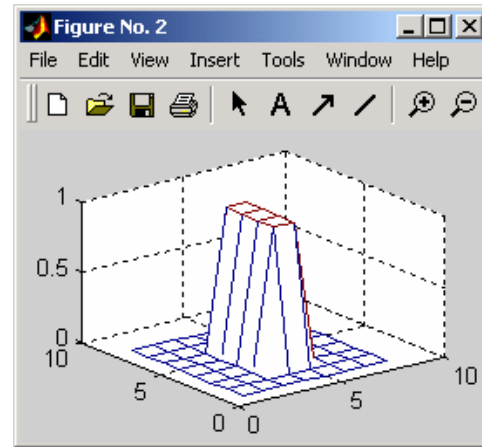


Figure 6.3: Mesh Plot of Original Image

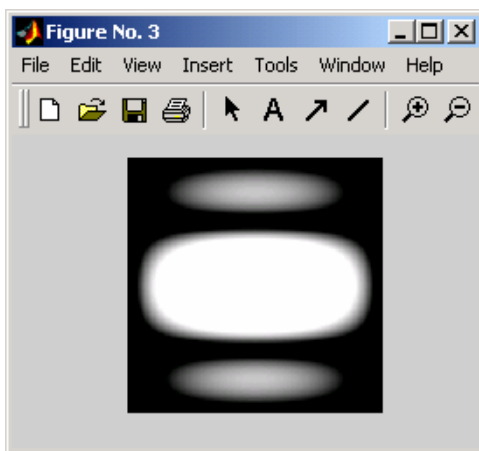


Figure 6.4: 2-D FFT of Original Image

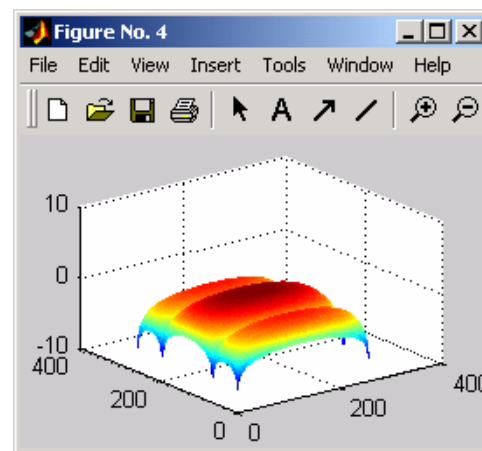


Figure 6.5: Mesh Plot of 2-D FFT

6.2. Convolution – Convolution is a linear filtering method commonly used in image processing. Convolution is the algebraic process of multiplying two polynomials. An image is an array of polynomials whose pixel values represent the coefficients of the polynomials. Therefore, two images can be multiplied together to produce a new image through the process of convolution. If the convolution kernel, or filter, is large, this can be a very tedious process involving many multiplication steps. However, the convolution theorem states that convolution is the same as the inverse Fourier Transform of the multiplication of two Fourier Transforms. In MATLAB, **conv2** is used to perform a two-dimensional convolution of two matrices. This can also be accomplished by taking the **ifft2** of the multiplication of two **fft2**'s. When this is done, though, both matrices' dimensions must be the same. This is not required when using **conv2**. Convolution is a neighborhood operation, since it uses the values of neighboring pixels in determining what the new pixel value will be. When MATLAB performs a convolution, it rotates the convolution kernel by 180° and multiplies it with a selected area on the original image, centered about a specific pixel. This pixel takes on the value of the sum of each original pixel value multiplied with its corresponding pixel value in the

convolution kernel. Then the kernel slides to the next pixel and the process is repeated, until all pixel values have been changed. If a 3x3 kernel is convolved with an image, each pixel will take on a new value related to the sum of itself, multiplied by the center of the convolution kernel, and its eight neighboring pixels multiplied by their own corresponding pixel value in the kernel. Example 6.2, below, for a demonstration of convolution.

Example 6.2

This example demonstrates that the convolution of two images is the same as inverse Fourier Transform of the multiplication of the Fourier Transforms of the two images. The M-file in Figure 6.6 contains the code necessary to demonstrate this task.

```

1 % Tutorial M-file
2 % Created by: Someone
3 % Created on: 9/11/03
4 % Last revised: 10/7/03
5
6 A = [0 0 0 0 0 0 0 0;      %8x8 matrix for testing convolution
7      0 0 0 0 0 0 0 0;
8      0 0 0 1 1 0 0 0;
9      0 0 0 1 1 0 0 0;
10     0 0 0 1 1 0 0 0;
11     0 0 0 1 1 0 0 0;
12     0 0 0 0 0 0 0 0;
13     0 0 0 0 0 0 0 0];
14
15 k = [1/9 1/9 1/9;          %3x3 test kernel for convolution
16     1/9 1/9 1/9;
17     1/9 1/9 1/9];
18
19 A(10,10) = 0;              %zero pad A to be size 10x10
20 B = conv2(A,k);            %convolve A with k
21 subplot(3,1,1); mesh(A);title('A');    %plot mesh image of A in 3 row by 1 column image, space 1
22 subplot(3,1,2); mesh(k);title('k');    %plot mesh image of k in 3 row by 1 column image, space 2
23 subplot(3,1,3); mesh(B);title('Convolution 1'); %plot mesh image of Convolution in 3 row by
24                                         %1 column image, space 3
25
26 k(10,10) = 0;              %zero pad k to be same size as A
27 C = ifft2(fft2(A).*fft2(k)); %comput conv. as IFFT of mult of 2 FFT's
28 figure;mesh(C);title('Convolution 2'); %display mesh plot of Conv. of A and k

```

Figure 6.6: M-file for Convolution

The “image” is the same as that used in Example 6.1. The convolution kernel is a 3x3 matrix with all values the same and scaled to the size of the matrix. This type of kernel, as you will see, has a low pass characteristic that tends to smooth out high frequency content in the original image. The plots that were created by this M-file are all displayed as mesh plots so that it is easier to view what effect the convolution kernel has on the original image. New to this example is the use of **subplot**. This function allows the user to place more than one plot in a figure window. In this case, there are three images in the figure window. Figure 6.7, below, depicts the original image (“A”), the convolution kernel

(“k”), and the result of the convolution of these two matrices (“Convolution 1”). Figure 6.8, below, is an image of the inverse two-dimensional FFT of the multiplication of the two-dimensional FFT’s of the two matrices. Notice how both methods provide the same results. The low pass characteristics of the convolution kernel are evident in the result. The peak has been eroded away and is now not as intense as before.

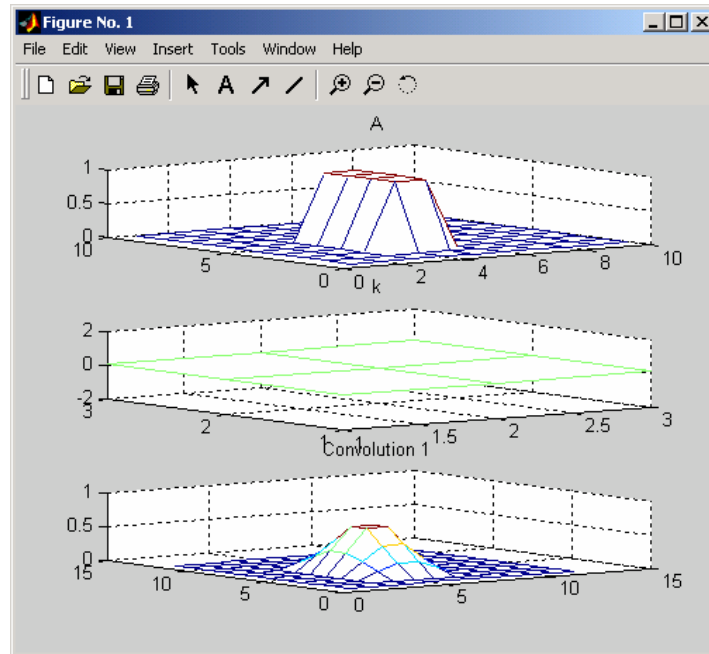


Figure 6.7: Convolution using conv2(A,k)

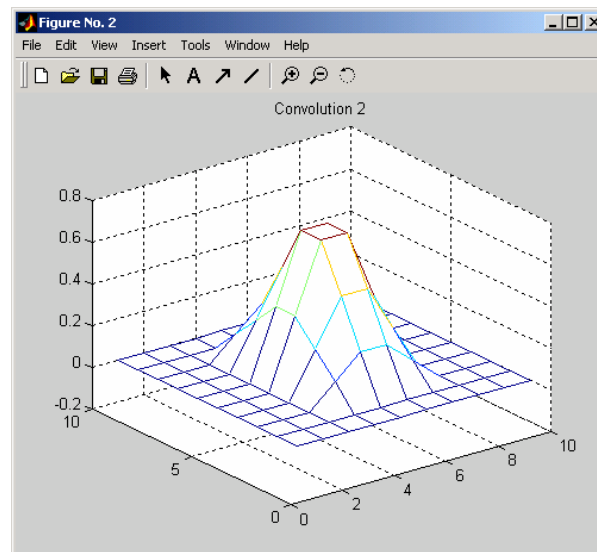


Figure 6.8: Convolution Using ifft2(fft2(A).*fft2(k))

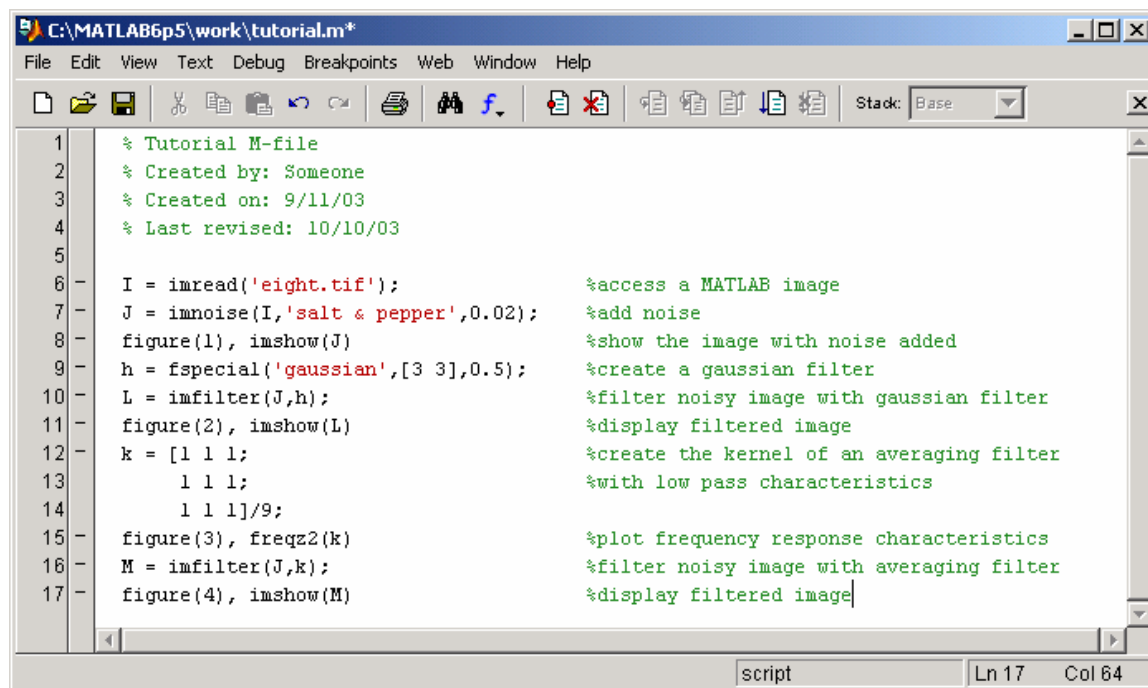
7. Filters

7.1. Filters – Image processing is based on filtering the content of images. Filtering is used to modify an image in some way. This could entail blurring, deblurring, locating certain features within an image, etc... Linear filtering is accomplished using convolution, as discussed above. A filter, or convolution kernel as it is also known, is basically an algorithm for modifying a pixel value, given the original value of the pixel and the values of the pixels surrounding it. There are literally hundreds of types of filters that are used in image processing. However, we will concentrate on several common ones.

7.2. Low Pass Filters – The first filters we will talk about are low pass filters. These filters blur high frequency areas of images. This can sometimes be useful when attempting to remove unwanted noise from an image. However, these filters do not discriminate between noise and edges, so they tend to smooth out content that should not be smoothed out. Example 6.2, above, provides an example of a basic low pass filter. The convolution kernel values can be modified to achieve desired low pass filter characteristics. See Example 7.1, below, on how to load an image and then apply a low pass filter to it.

Example 7.1

This example demonstrates how to load an image that is stored in MATLAB's files, and how to filter the content of the image. The same image is filtered by two different low pass filters. The goal is to remove the noise present in the image. The M-File in Figure 7.1, below, contains the code for this example. The image, eight.tif, is a MATLAB example image.



```
1 % Tutorial M-file
2 % Created by: Someone
3 % Created on: 9/11/03
4 % Last revised: 10/10/03
5
6 I = imread('eight.tif'); %access a MATLAB image
7 J = imnoise(I,'salt & pepper',0.02); %add noise
8 figure(1), imshow(J) %show the image with noise added
9 h = fspecial('gaussian',[3 3],0.5); %create a gaussian filter
10 L = imfilter(J,h); %filter noisy image with gaussian filter
11 figure(2), imshow(L) %display filtered image
12 k = [1 1 1; %create the kernel of an averaging filter
13     1 1 1; %with low pass characteristics
14     1 1 1]/9;
15 figure(3), freqz2(k) %plot frequency response characteristics
16 M = imfilter(J,k); %filter noisy image with averaging filter
17 figure(4), imshow(M) %display filtered image
```

Figure 7.1: M-file for Low Pass Filter Design

The images generated by the M-file in Figure 7.1 are pictured in Figures 7.2-7.5. Figure 7.2 is a MATLAB image with salt and pepper noise added to it. Figure 7.3 is the result of a 3x3 Gaussian filter with low pass characteristics applied to the image in Figure 7.2. Figure 7.4 is the frequency response of a 3x3 averaging filter with all values equal and scaled to the size of the filter. Notice the low pass characteristics of this filter. Figure 7.5 is the result of the filter depicted in Figure 7.4 applied to the image in Figure 7.2.

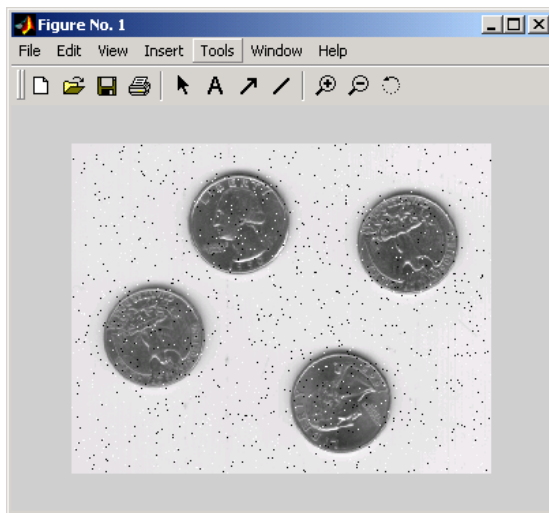


Figure 7.2: Noisy Image

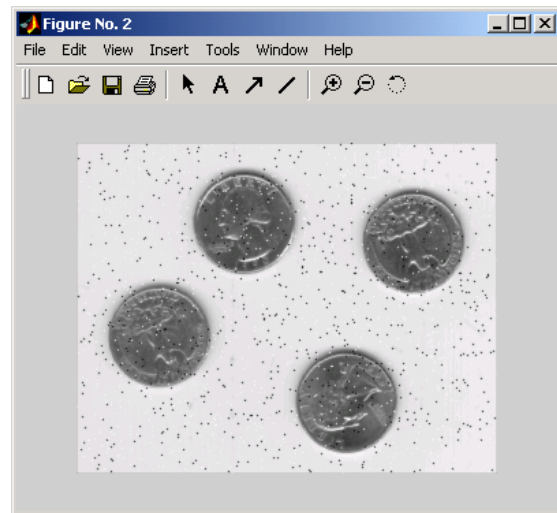


Figure 7.3: Gaussian Filtered Image

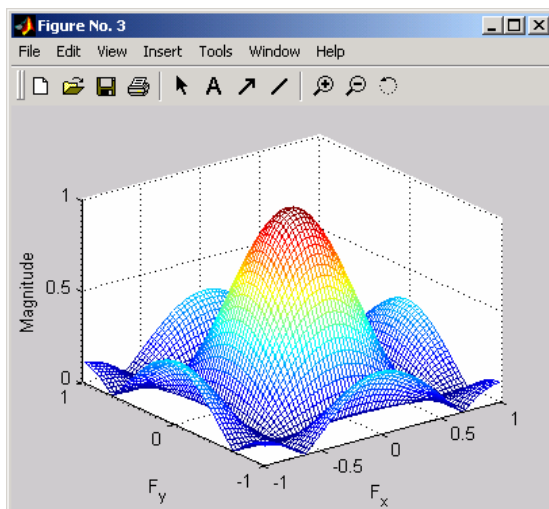


Figure 7.4: Averaging Filter Response

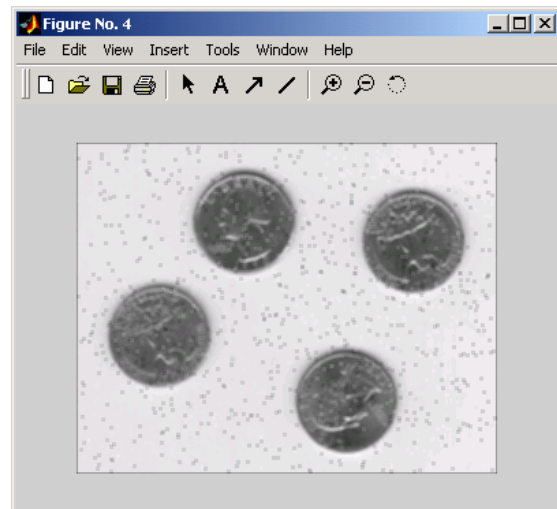


Figure 7.5: Averaging Filtered Image

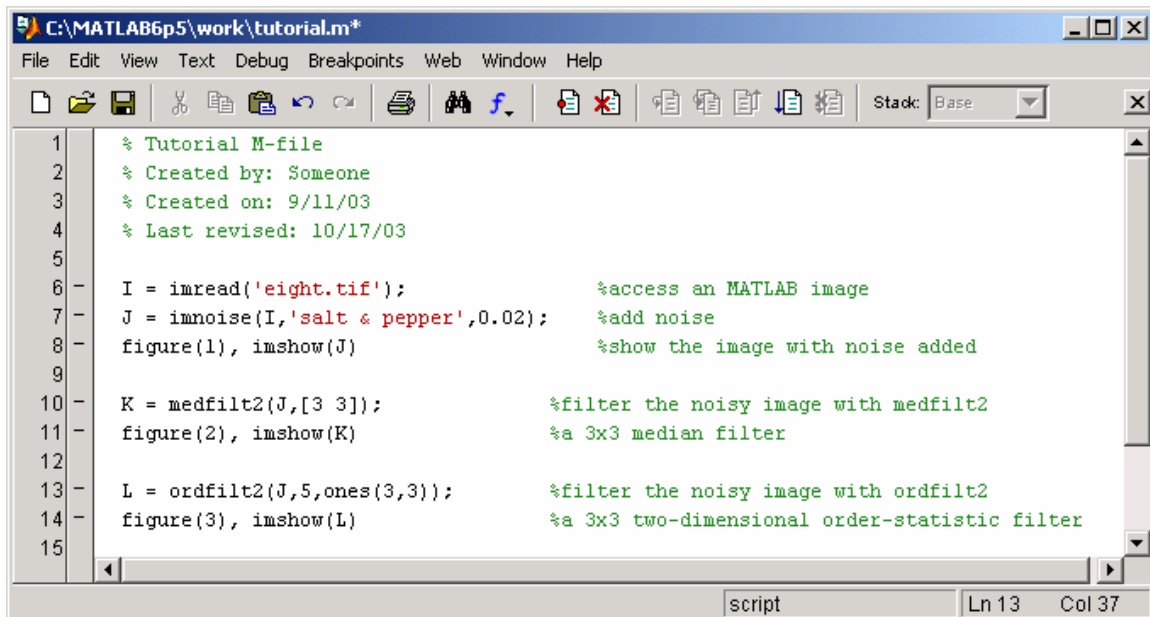
As you can see some of the noise apparent in the image in Figure 7.2 has been blurred by both filters. However, neither does a good job removing the noise. In fact, if the noise was to be adequately attenuated, the coins in the images would become so blurred, the filtered image would be much worse than the original image. Low pass filters are pretty good at removing noise with pixel values close to the surrounding pixel values. However, this is not always the

case. Fortunately, low pass filters are not the only filters capable of removing noise.

- 7.3. Median Filters** – Median Filters can be very useful for removing noise from images. A median filter is like an averaging filter in some ways. The averaging filter examines the pixel in question and its neighbor's pixel values and returns the mean of these pixel values. The median filter looks at this same neighborhood of pixels, but returns the median value. In this way noise can be removed, but edges are not blurred as much, since the median filter is better at ignoring large discrepancies in pixel values. See Example 7.2, below, for how to perform a median filtering operation.

Example 7.2

This example uses two types of median filters that both output the same result. The first filter is **medfilt2**, which takes the median value of the pixel in question and its neighbors. In this case it outputs the median value of nine pixels being examined. The second filter, **ordfilt2**, does the exact same thing in this configuration, but can be configured to perform other types of filtering. In this case, it looks at every pixel in the 3x3 matrix and outputs the value in the fifth position of rank, which is the median position. In other words it outputs a value, where half the pixel values are greater and half are less, in the matrix.



```
1 % Tutorial M-file
2 % Created by: Someone
3 % Created on: 9/11/03
4 % Last revised: 10/17/03
5
6 I = imread('eight.tif');           %access an MATLAB image
7 J = imnoise(I,'salt & pepper',0.02); %add noise
8 figure(1), imshow(J)              %show the image with noise added
9
10 K = medfilt2(J,[3 3]);             %filter the noisy image with medfilt2
11 figure(2), imshow(K)              %a 3x3 median filter
12
13 L = ordfilt2(J,5,ones(3,3));       %filter the noisy image with ordfilt2
14 figure(3), imshow(L)              %a 3x3 two-dimensional order-statistic filter
15
```

Figure 7.6: M-file for Median Filter Design

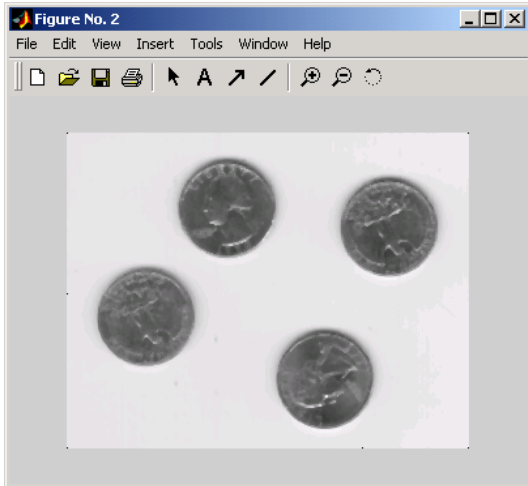


Figure 7.7: medfilt2

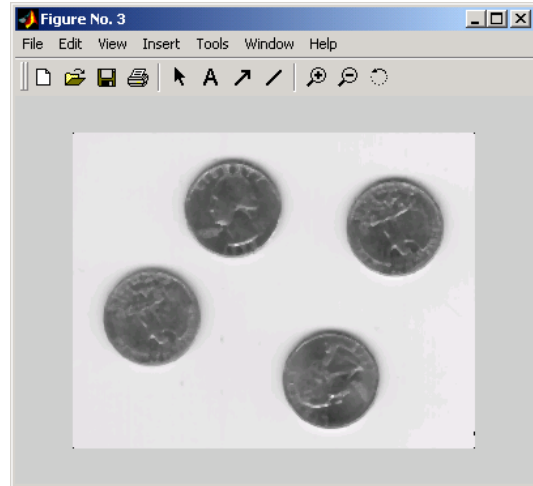


Figure 7.8: ordfilt2

Figure 7.6, above depicts the M-file used in this example. The original image in this example is the image in Figure 7.2. Figure 7.7, above, is the output of the image in Figure 7.2, filtered with a 3x3 two-dimensional median filter. Figure 7.8, above, is the same as Figure 7.7, but was achieved by filtering the image in Figure 7.2 with **ordfilt2**, configured to produce the same result as **medfilt2**. Notice how both filters produce the same result. Each is able to remove the noise, without blurring the edges in the image too much.

7.4 Erosion and Dilation – Erosion and Dilation are similar operations to median filtering in that they both are neighborhood operations. The erosion operation examines the value of a pixel and its neighbors and sets the output value equal to the minimum of the input pixel values. Dilation, on the other hand, examines the same pixels and outputs the maximum of these pixels. In MATLAB erosion and dilation can be accomplished by the **imerode** and **imdilate** functions, respectively, accompanied by the **strel** function. Example 7.3 below, demonstrates erosion and dilation.

Example 7.3

In order to erode or dilate an image you must first specify to what extent and in what way you would like to erode or dilate the image. This is accomplished by creating a structuring element by using the **strel** function. There are many types of structuring elements, each with their own unique properties. For this example, the square shape provides a 5x5 square structuring element. To find other shapes for structuring elements, look up **strel** in MATLAB's help. Figure 7.9 contains the M-file for this example. The image used in this example is the same image of quarters used in the previous two examples. Figure 7.10 depicts erosion of the original image, while Figure 7.11 contains a dilation of the original image. The intent of this example was to exaggerate the results of the erosion and dilation operations. As you can see in the eroded image, the quarters are very dark, while in the dilated image the quarters are especially bright. In actual applications the structuring element must be configured to process the image according to desired results.

```

1  % Tutorial M-file
2  % Created by: Someone
3  % Created on: 9/11/03
4  % Last revised: 12/10/03
5
6  A = imread('eight.tif');           %access an MATLAB image
7  se = strel('square',5);           %create 5x5 square structuring element
8  B = imerode(A,se);                %erode image
9  C = imdilate(A,se);               %dilate image
10 figure(1), imshow(B), title('Eroded') %display eroded image
11 figure(2), imshow(C), title('Dilated') %display dilated image

```

Figure 7.9: M-file for Erosion and Dilation

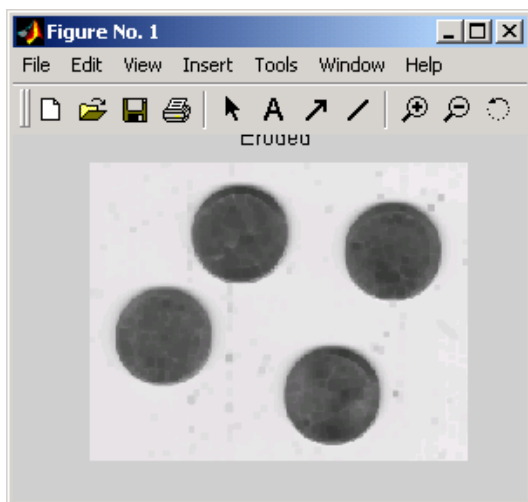


Figure 7.10: Erosion

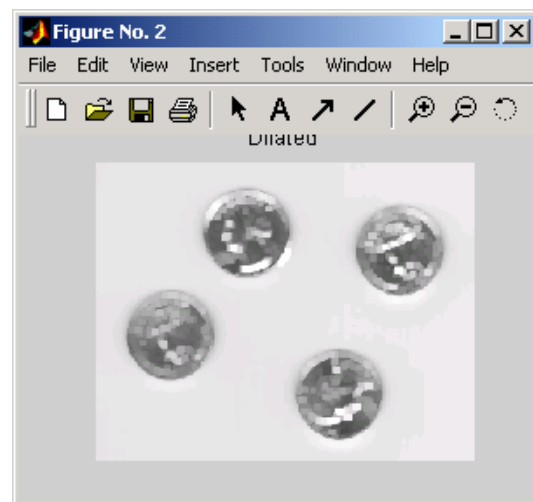


Figure 7.11: Dilation

7.5 Edge Detectors – Edge detectors are very useful for locating objects within images. There are many different kinds of edge detectors, but we will concentrate on two: the Sobel edge detector and the Canny edge detector. The Sobel edge detector is able to look for strong edges in the horizontal direction, vertical direction, or both directions. The Canny edge detector detects all strong edges plus it will find weak edges that are associated with strong edges. Both of these edge detectors return binary images with the edges shown in white on a black background. Example 7.4, below, demonstrates the use of these edge detectors.

Example 7.4

The Canny and Sobel edge detectors are both demonstrated in this example. Figure 7.12, below, is a sample M-file for performing these operations. The

image used is the MATLAB image, `rice.tif`, which can be found in the manner described in Example 4.1. Two methods for performing edge detection using the Sobel method are shown. The first method uses the MATLAB functions, **`fspecial`**, which creates the filter, and **`imfilter`**, which applies the filter to the image. The second method uses the MATLAB function, **`edge`**, in which you must specify the type of edge detection method desired. Sobel was used as the first edge detection method, while Canny was used as the next type. Figure 7.13, below, displays the results of the M-file in figure 7.12. The first image is the original image; the image denoted Horizontal Sobel is the result of using **`fspecial`** and **`imfilter`**. The image labeled Sobel is the result of using the edge filter with Sobel specified, while the image labeled Canny has Canny specified.

```

1 % Tutorial M-file
2 % Created by: Someone
3 % Created on: 9/11/03
4 % Last revised: 11/12/03
5
6 I = imread('rice.tif'); %load MATLAB image
7 subplot(2,2,1);imshow(I);title('Original Image');
8 h = fspecial('sobel'); %create Sobel (horizontal) edge detector
9 A = imfilter(I,h); %filter rice image with Sobel filter
10 %to get vertical filter transpose h => use h'
11 subplot(2,2,2);imshow(A);title('Horizontal Sobel');
12 B = edge(I,'sobel',[],'both'); %filter image with Sobel edge detector
13 %[] is for threshold value, 'both' is for horizontal and vertical
14 %edge detection ('horizontal' or 'vertical' may be specified, also).
15 subplot(2,2,3);imshow(B);title('Sobel');
16 C = edge(I,'canny',[],1); %filter image with Canny edge detector
17 %[] is for threshold value and 1 is for standard deviation of Gaussian
18 %filter
19 subplot(2,2,4);imshow(C);title('Canny');

```

Figure 7.12: M-File for Edge Detection

The Zoom In tool was used to depict the detail in the images more clearly. As you can see, the filter used to create the Horizontal Sobel image detects horizontal edges much more readily than vertical edges. The filter used to create the Sobel image detected both horizontal and vertical edges. This resulted from MATLAB looking for both horizontal and vertical edges independently and then summing them. The Canny image demonstrates how well the Canny method detects all edges. The Canny method does not only look for strong edges, as in the Sobel method, but also will look for weak edges that are connected to strong edges and show those, too.

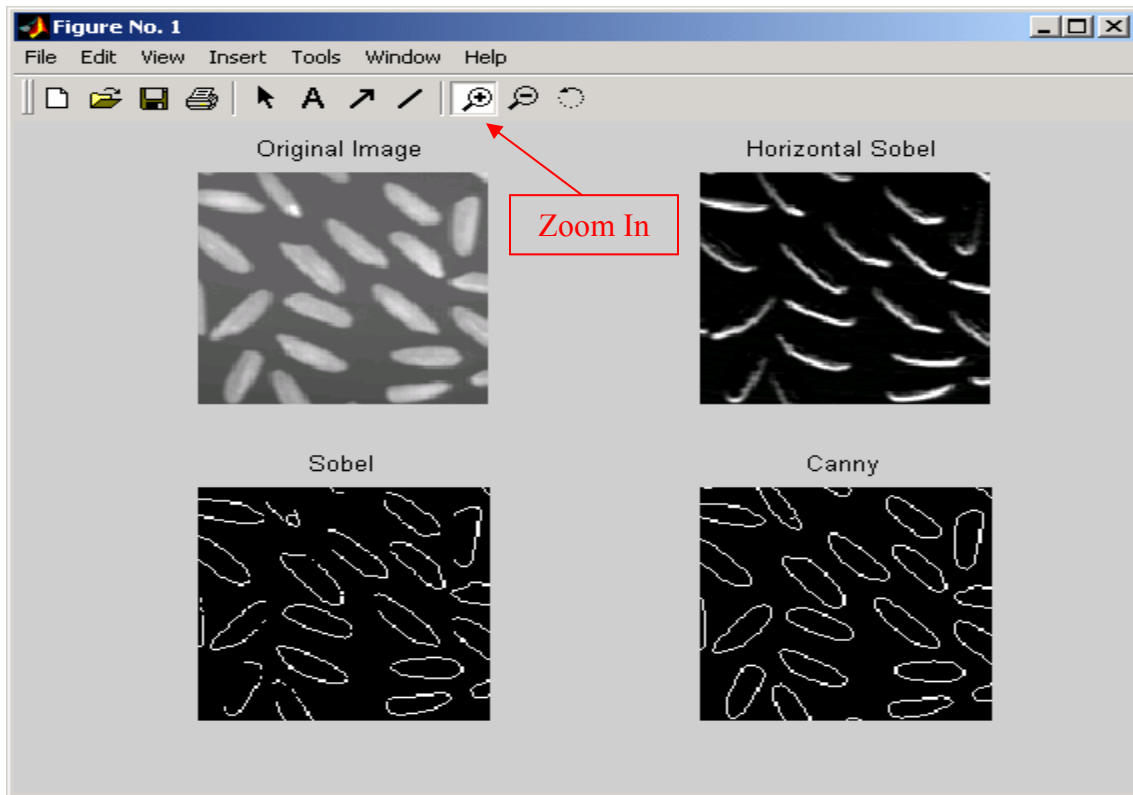


Figure 7.13: Images Created by Different Edge Detection Methods

- 7.6. Segmentation** – Segmentation is the process of fractioning an image into its component objects. This can be accomplished in various ways in MATLAB. One way is to use a combination of morphological operations to segment touching objects within an image. This is illustrated in Example 7.5. Another method is to use a combination of dilation and erosion to segment objects. The MATLAB function **bwperim** performs this operation on binary images.

Example 7.5

This example demonstrates the process of Watershed segmentation. The M-file for this demonstration is pictured in Figure 7.14. The first step in this process is to load an image in the way described in Example 7.1. The next step is to create a structural element, using **strel**, that resembles the objects present in the image. In this case the structural element shape “diamond” was used, however, “line” would have also worked. The third step is to perform both top-hat and bottom-hat filtering on the image according to the structuring element using **imtophat** and **imbothat**, respectively. Top-hat filtering is used to intensify valleys in an image, while bottom-hat filtering enhances contrast. The next step is to enhance the contrast even more by combining the products of the top-hat and bottom-hat filtering using **imsubtract** and **imadd**, which subtract and add images, respectively. The next step is to complement the enhanced image, using **imcomplement**, which is the same as creating the negative of the object. The next step is to use both the **imextendedmin** function and the **imimposemin** function to create larger valleys and set the valley value to the minimum

possible for the class size, or zero. The last step is to perform watershed segmentation on the product of these morphological operations, using the **watershed** function. In order to view the watershed segments better the image was then converted to RGB, using the **label2rgb** function, enabling each labeled watershed segment as a different color. Figure 7.15, below, contains the original image and the segmented image.

```

C:\MATLAB6p5\work\tutorial.m*
File Edit View Text Debug Breakpoints Web Window Help
[Icons] Stack: Base
1  %% Tutorial M-file
2  % Created by: Someone
3  % Created on: 9/11/03
4  % Last revised: 12/12/03
5
6  A = imread('testpatl.tif');
7  subplot(1,2,1); imshow(A), title('Original Image'); %original image
8  se = strel('diamond',7); %structuring element
9  Itop = imtophat(A, se); %top hat of image
10 Ibot = imbothat(A, se); %bottom hat of image
11 Ienhance = imsubtract(imadd(Itop, afm), Ibot); %enhance contrast between objects
12 Iec = imcomplement(Ienhance); %complement - highlight intensity valleys
13 Iemin = imextendedmin(Iec, 22); %extend minima of complement
14 Iimpose = imimposemin(Iec, Iemin); %impose minima on complement
15 wat = watershed(Iimpose); %watershed segmentation
16 rgb = label2rgb(wat); %convert to RGB to highlight contrast better
17 subplot(1,2,2); imshow(rgb); title('Watershed Segmented Image'); %display watershed
script Ln 17 Col 84

```

Figure 7.14: M-file for Watershed Segmentation

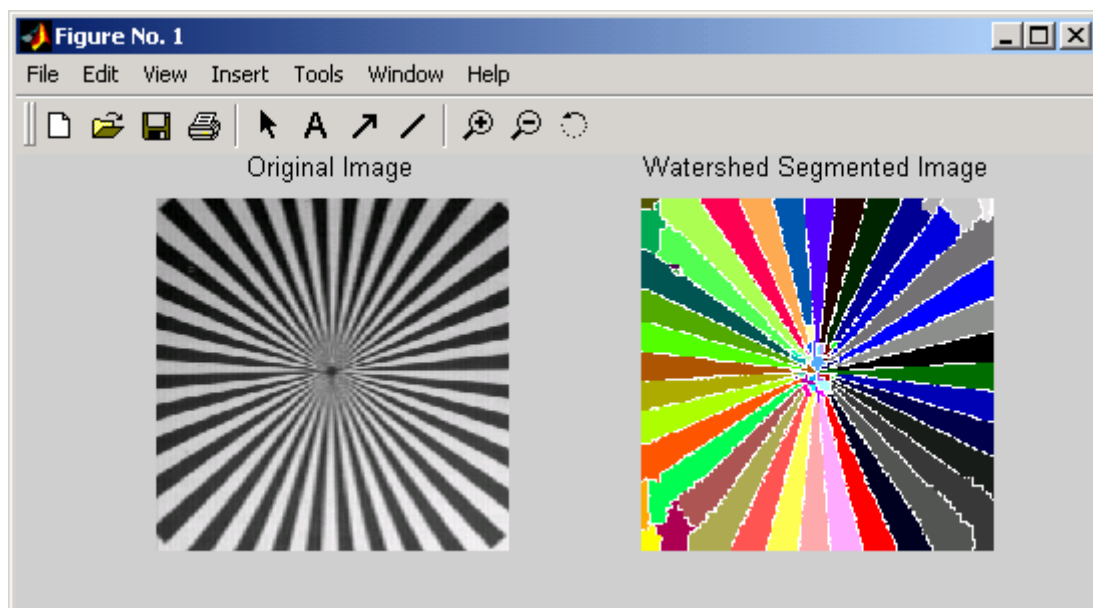


Figure 7.15: Original Image and Segmented Image